

# About Jamboree

Jamboree has helped thousands of students like you make it to top colleges abroad. Be it GMAT, GRE or SAT, their unique problem-solving methods ensure maximum scores with minimum effort. They recently launched a feature where students/learners can come to their website and check their probability of getting into the IVY league college. This feature estimates the chances of graduate admission from an Indian perspective.

## Business Problem

To help Jamboree understand what factors are important in graduate admissions and how these factors are interrelated among themselves. It will also help predict one's chances of admission given the rest of the variables.

```
In [1]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import r2_score

from statsmodels.stats.outliers_influence import variance_inflation_factor
from scipy import stats
```

```
In [2]: # Loading dataset
df = pd.read_csv("/Users/bose/Downloads/Jamboree.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	1	337	118	4	4.5	4.5	9.65	1	0.92
1	2	324	107	4	4.0	4.5	8.87	1	0.76
2	3	316	104	3	3.0	3.5	8.00	1	0.72
3	4	322	110	3	3.5	2.5	8.67	1	0.80
4	5	314	103	2	2.0	3.0	8.21	0	0.65

```
In [4]: df.tail()
```

Out [4]:

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
<b>495</b>	496	332	108	5	4.5	4.0	9.02	1	0.87
<b>496</b>	497	337	117	5	5.0	5.0	9.87	1	0.96
<b>497</b>	498	330	120	5	4.5	5.0	9.56	1	0.93
<b>498</b>	499	312	103	4	4.0	5.0	8.43	0	0.73
<b>499</b>	500	327	113	4	4.5	4.5	9.04	0	0.84

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Serial No.            500 non-null   int64
1   GRE Score              500 non-null   int64
2   TOEFL Score            500 non-null   int64
3   University Rating      500 non-null   int64
4   SOP                    500 non-null   float64
5   LOR                    500 non-null   float64
6   CGPA                   500 non-null   float64
7   Research               500 non-null   int64
8   Chance of Admit        500 non-null   float64
dtypes: float64(4), int64(5)
memory usage: 35.3 KB
```

In [6]: `df.shape`

Out[6]: (500, 9)

In [7]: `df.dtypes`

```
Out[7]: Serial No.            int64
GRE Score              int64
TOEFL Score            int64
University Rating      int64
SOP                    float64
LOR                    float64
CGPA                   float64
Research               int64
Chance of Admit        float64
dtype: object
```

In [8]: `df.describe()`

Out [8]:

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGI
<b>count</b>	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000
<b>mean</b>	250.500000	316.472000	107.192000	3.114000	3.374000	3.484000	8.576400
<b>std</b>	144.481833	11.295148	6.081868	1.143512	0.991004	0.925450	0.604800
<b>min</b>	1.000000	290.000000	92.000000	1.000000	1.000000	1.000000	6.800000
<b>25%</b>	125.750000	308.000000	103.000000	2.000000	2.500000	3.000000	8.127500
<b>50%</b>	250.500000	317.000000	107.000000	3.000000	3.500000	3.500000	8.560000
<b>75%</b>	375.250000	325.000000	112.000000	4.000000	4.000000	4.000000	9.040000
<b>max</b>	500.000000	340.000000	120.000000	5.000000	5.000000	5.000000	9.920000

In [9]: *#Checking for Null values -*  
`df.isnull().sum()`

Out [9]:

Serial No.	0
GRE Score	0
TOEFL Score	0
University Rating	0
SOP	0
LOR	0
CGPA	0
Research	0
Chance of Admit	0

dtype: int64

There are no null values in the dataset

In [10]: *#Removing column Serial No. -*  
`df.drop(columns=['Serial No.'], inplace=True)`

In [11]: *#Checking for duplicates*  
`df.duplicated().sum()`

Out [11]: 0

There are no duplicates in the given dataset

In [12]:

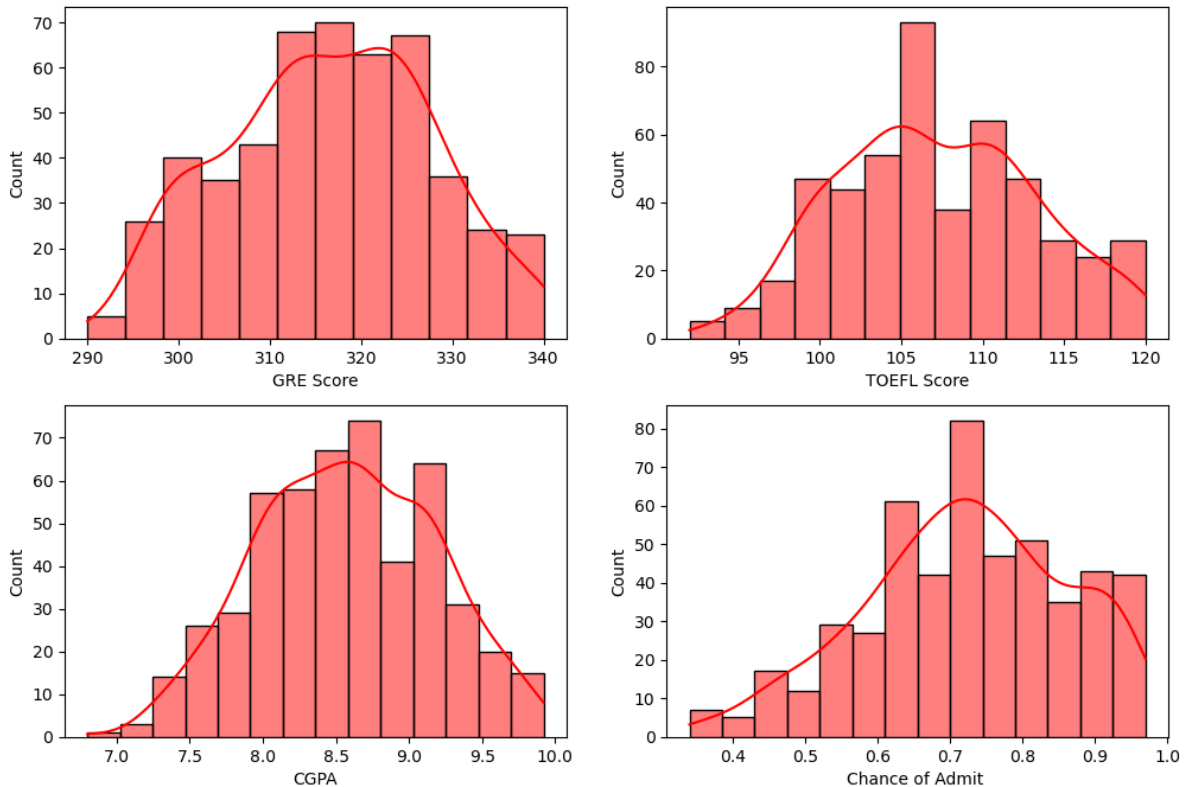
```
cat_cols = ['University Rating', 'SOP', 'LOR ', 'Research']
num_cols = ['GRE Score', 'TOEFL Score', 'CGPA']
target = 'Chance of Admit '
```

## Univariate Analysis

In [13]: *#Distribution of Continuous numerical features*

```
rows, cols = 2, 2
fig, axs = plt.subplots(rows, cols, figsize=(12, 8))
index = 0
for row in range(rows):
    for col in range(cols):
        sns.histplot(df[num_cols[index]], kde=True, ax=axs[row, col], color=
            index += 1
    break
```

```
sns.histplot(df[num_cols[-1]], kde=True, ax=axes[1,0], color='r')
sns.histplot(df[target], kde=True, ax=axes[1,1], color='r')
plt.show()
```




### Observation -

- Most students have GRE score in the range of (310-330)
- Most students have TOEFL score in the range (105-110)
- Majority of students have CGPA in the range of (8.0 - 9.0)
- Average Chance of Admit for the student comes out to 0.72

```
In [14]: #Checking for outliers using boxplots
rows, cols = 2, 2
fig, axes = plt.subplots(rows, cols, figsize=(12, 7))

index = 0
for col in range(cols):
    sns.boxplot(x=num_cols[index], data=df, ax=axes[0,index],color='pink')
    index += 1

sns.boxplot(x=num_cols[-1], data=df, ax=axes[1,0],color='pink')
sns.boxplot(x=target, data=df, ax=axes[1,1],color='pink')
plt.show()
```

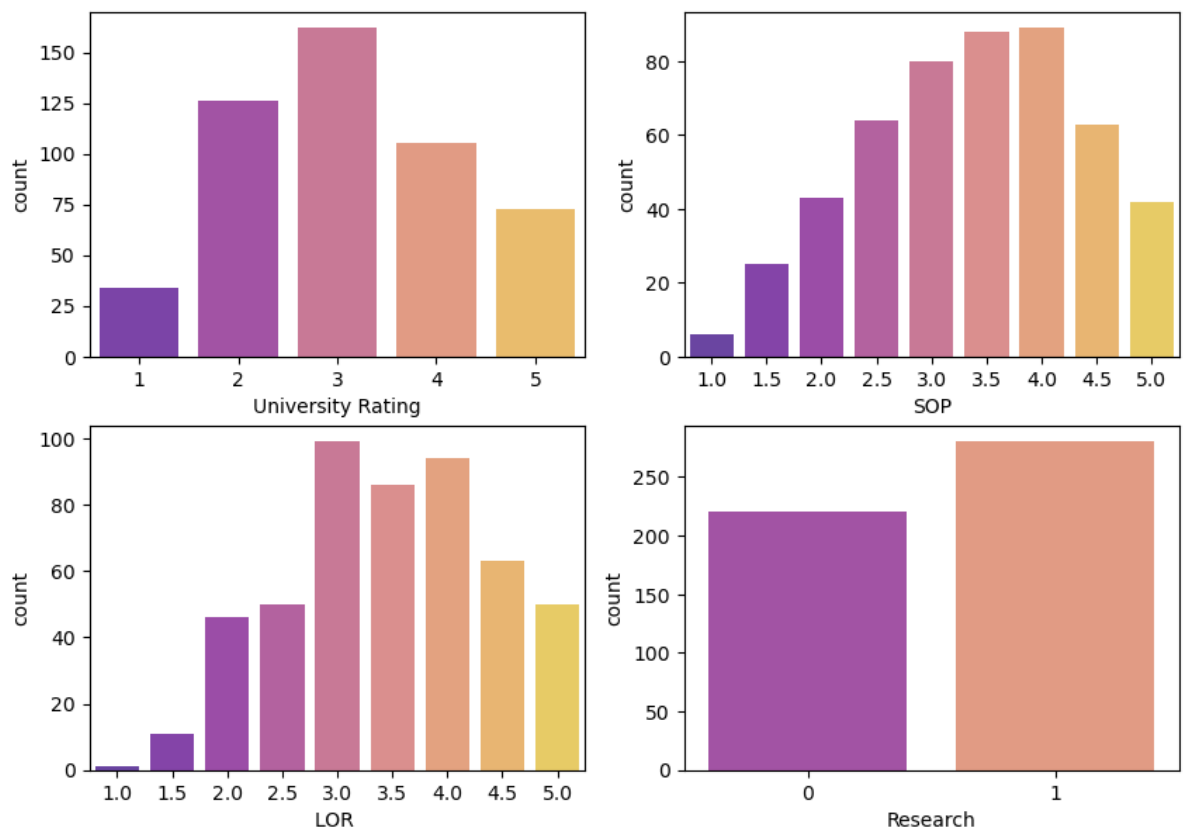


There are no outliers in the dataset

```
In [15]: #Countplots for categorical variables
cols, rows = 2, 2
fig, axs = plt.subplots(rows, cols, figsize=(10, 7))

index = 0
for row in range(rows):
    for col in range(cols):
        sns.countplot(x=cat_cols[index], data=df, ax=axs[row, col], alpha=0.5)
        index += 1

plt.show()
```



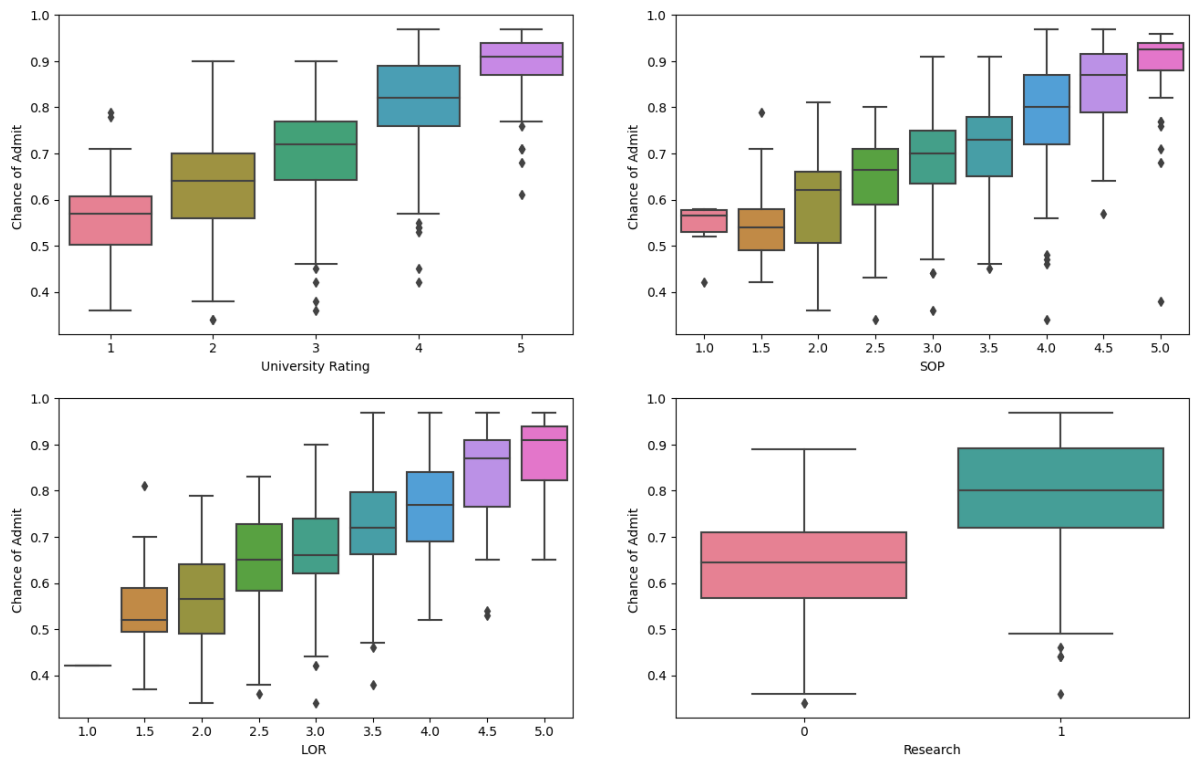
### Observation -

- Most no of students have a University Ranking of 3
- SOP score of 3.5 and 4.0 is most common among students
- Highest number of students have an LOR score of 3.0
- No of students who have conducted research is more than the no of students who have not

## Bivariate Analysis

```
In [16]: #Boxplot to represent the effect of categorical values on Chance of Admit
rows, cols = 2,2
fig, axs = plt.subplots(rows, cols, figsize=(16,10))


index = 0
for row in range(rows):
    for col in range(cols):
        sns.boxplot(x=cat_cols[index], y=target, data=df, ax=axs[row,col],
                    index += 1
```



### Observation -

- As the University Rating increases, Chance of Admit also increases
- Same is the case with SOP and LOR. A Candidate with higher rating of SOP and LOR have a higher Chance of Admit
- A student with more research experience have higher Chance of Admit than a student without research experience


```
In [17]: sns.pairplot(df[num_cols])
plt.show()
```



The numerical columns seems to have a correlation with each other

```
In [18]: plt.figure(figsize=(10,8))  
sns.heatmap(df.corr(), annot=True)  
plt.show()
```





### Observation -

- **Chance of Admit** have a **high correlation** with **GRE Score, TOEFL Score and CGPA**
- Research seems to have low correlation with other features

## Model Building

```
In [19]: X = df.drop(columns=[target])
         y = df[target]
```

```
In [20]: #Standardize the dataset
         sc = StandardScaler()
         X = sc.fit_transform(X)
```

```
In [21]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ran
```

```
In [22]: print(X_train.shape, y_train.shape)
         print(X_test.shape, y_test.shape)
```

```
(350, 7) (350,)
(150, 7) (150,)
```

```
In [23]: def adjusted_r2(r2, p, n):
         .....
```

```
         n: no of samples
```

```

p: no of predictors
r2: r2 score
"""
adj_r2 = 1 - ((1-r2)*(n-1) / (n-p-1))
return adj_r2

def get_metrics(y_true, y_pred, p=None):
    n = y_true.shape[0]
    mse = np.sum((y_true - y_pred)**2) / n
    rmse = np.sqrt(mse)
    mae = np.mean(np.abs(y_true - y_pred))
    score = r2_score(y_true, y_pred)
    adj_r2 = None
    if p is not None:
        adj_r2 = adjusted_r2(score, p, n)

    res = {
        "mean_absolute_error": round(mae, 2),
        "rmse": round(rmse, 2),
        "r2_score": round(score, 2),
        "adj_r2": round(adj_r2, 2)
    }
    return res

```

```

In [24]: def train_model(X_train, y_train, X_test, y_test, cols, model_name="linear",
                        model = None):
    if model_name == "lasso":
        model = Lasso(alpha=alpha)
    elif model_name == "ridge":
        model = Ridge(alpha=alpha)
    else:
        model = LinearRegression()

    model.fit(X_train, y_train)
    y_pred_train = model.predict(X_train)
    y_pred_test = model.predict(X_test)
    p = X_train.shape[1]
    train_res = get_metrics(y_train, y_pred_train, p)
    test_res = get_metrics(y_test, y_pred_test, p)

    print(f"\n----- {model_name.title()} Regression Model -----\n")
    print(f"Train MAE: {train_res['mean_absolute_error']} Test MAE: {test_res['mean_absolute_error']}")
    print(f"Train RMSE: {train_res['rmse']} Test RMSE: {test_res['rmse']}")
    print(f"Train R2_score: {train_res['r2_score']} Test R2_score: {test_res['r2_score']}")
    print(f"Train Adjusted_R2: {train_res['adj_r2']} Test Adjusted_R2: {test_res['adj_r2']}")
    print(f"Intercept: {model.intercept_}")
    coef_df = pd.DataFrame({"Column": cols, "Coefficient": model.coef_})
    print(coef_df)
    print("-"*50)
    return model

```

```

In [25]: train_model(X_train, y_train, X_test, y_test, df.columns[:-1], "linear")
         train_model(X_train, y_train, X_test, y_test, df.columns[:-1], "ridge")
         train_model(X_train, y_train, X_test, y_test, df.columns[:-1], "lasso", 0.001)

```

----- Linear Regression Model -----

Train MAE: 0.04 Test MAE: 0.04  
 Train RMSE: 0.06 Test RMSE: 0.06  
 Train R2\_score: 0.82 Test R2\_score: 0.82  
 Train Adjusted\_R2: 0.82 Test Adjusted\_R2: 0.81  
 Intercept: 0.724978121476996

	Column	Coefficient
0	GRE Score	0.018657
1	TOEFL Score	0.023176
2	University Rating	0.011565
3	SOP	-0.000999
4	LOR	0.012497
5	CGPA	0.064671
6	Research	0.013968

-----

----- Ridge Regression Model -----

Train MAE: 0.04 Test MAE: 0.04  
 Train RMSE: 0.06 Test RMSE: 0.06  
 Train R2\_score: 0.82 Test R2\_score: 0.82  
 Train Adjusted\_R2: 0.82 Test Adjusted\_R2: 0.81  
 Intercept: 0.7249823645841696

	Column	Coefficient
0	GRE Score	0.018902
1	TOEFL Score	0.023252
2	University Rating	0.011594
3	SOP	-0.000798
4	LOR	0.012539
5	CGPA	0.064004
6	Research	0.013990

-----

----- Lasso Regression Model -----

Train MAE: 0.04 Test MAE: 0.04  
 Train RMSE: 0.06 Test RMSE: 0.06  
 Train R2\_score: 0.82 Test R2\_score: 0.82  
 Train Adjusted\_R2: 0.82 Test Adjusted\_R2: 0.81  
 Intercept: 0.7249659139557142

	Column	Coefficient
0	GRE Score	0.018671
1	TOEFL Score	0.022770
2	University Rating	0.010909
3	SOP	0.000000
4	LOR	0.011752
5	CGPA	0.064483
6	Research	0.013401

-----

Out[25]: Lasso(alpha=0.001)

### Observation -

- Since there is no difference between train and test scores, we can say that there is no overfitting.
- There are no unnecessary independent variables in the data, as the value of R2 and Adjusted\_R2 are almost same.

## Assumptions of Linear Regression Model

Multicollinearity Check

```
In [26]: def vif(newdf):
# VIF dataframe
vif_data = pd.DataFrame()
vif_data["feature"] = newdf.columns

# calculating VIF for each feature
vif_data["VIF"] = [variance_inflation_factor(newdf.values, i)
                    for i in range(len(newdf.columns))]

return vif_data
```

```
In [27]: res = vif(df.iloc[:, :-1])
res
```

Out[27]:

	feature	VIF
0	GRE Score	1308.061089
1	TOEFL Score	1215.951898
2	University Rating	20.933361
3	SOP	35.265006
4	LOR	30.911476
5	CGPA	950.817985
6	Research	2.869493

```
In [28]: #Drop GRE Score and again calculate the VIF
res = vif(df.iloc[:, 1:-1])
res
```

Out[28]:

	feature	VIF
0	TOEFL Score	639.741892
1	University Rating	19.884298
2	SOP	33.733613
3	LOR	30.631503
4	CGPA	728.778312
5	Research	2.863301

```
In [29]: #Drop TOEFL Score and again calculate the VIF
res = vif(df.iloc[:, 2:-1])
res
```

Out[29]:

	feature	VIF
0	University Rating	19.777410
1	SOP	33.625178
2	LOR	30.356252
3	CGPA	25.101796
4	Research	2.842227

```
In [30]: #Drop SOP and again calculate VIF
res = vif(df.iloc[:,2:-1].drop(columns=['SOP']))
res
```

```
Out[30]:
```

	feature	VIF
0	University Rating	15.140770
1	LOR	26.918495
2	CGPA	22.369655
3	Research	2.819171

```
In [31]: #Drop LOR and again calculate VIF
newdf = df.iloc[:,2:-1].drop(columns=['SOP'])
newdf = newdf.drop(columns=['LOR'], axis=1)
res = vif(newdf)
res
```

```
Out[31]:
```

	feature	VIF
0	University Rating	12.498400
1	CGPA	11.040746
2	Research	2.783179

```
In [32]: #Drop University Rating and again calculate VIF
newdf = newdf.drop(columns=['University Rating'])
res = vif(newdf)
res
```

```
Out[32]:
```

	feature	VIF
0	CGPA	2.455008
1	Research	2.455008

```
In [33]: #Again train the model with only these two features
X = df[['CGPA', 'Research']]
sc = StandardScaler()
X = sc.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ran
```

```
In [34]: model = train_model(X_train, y_train, X_test, y_test, ['CGPA', 'Research'],
train_model(X_train, y_train, X_test, y_test, ['CGPA', 'Research'], "ridge")
train_model(X_train, y_train, X_test, y_test, ['CGPA', 'Research'], "lasso")
```

----- Linear Regression Model -----

```
Train MAE: 0.05 Test MAE: 0.05
Train RMSE: 0.06 Test RMSE: 0.07
Train R2_score: 0.78 Test R2_score: 0.81
Train Adjusted_R2: 0.78 Test Adjusted_R2: 0.81
Intercept: 0.7247774222727991
      Column Coefficient
0      CGPA      0.112050
1 Research      0.020205
```

-----

----- Ridge Regression Model -----

```
Train MAE: 0.05 Test MAE: 0.05
Train RMSE: 0.06 Test RMSE: 0.07
Train R2_score: 0.78 Test R2_score: 0.81
Train Adjusted_R2: 0.78 Test Adjusted_R2: 0.81
Intercept: 0.7247830300095277
      Column Coefficient
0      CGPA      0.111630
1 Research      0.020362
```

-----

----- Lasso Regression Model -----

```
Train MAE: 0.05 Test MAE: 0.05
Train RMSE: 0.06 Test RMSE: 0.07
Train R2_score: 0.78 Test R2_score: 0.81
Train Adjusted_R2: 0.78 Test Adjusted_R2: 0.81
Intercept: 0.7247713356661623
      Column Coefficient
0      CGPA      0.111344
1 Research      0.019571
```

-----


Out[34]: Lasso(alpha=0.001)

### Observation -

- The R2 score and Adjusted\_R2 score are almost same, even after removing collinear features using VIF and using only two features.
- Mean of Residuals : From the RMSE score we can say that the Mean of Residuals is nearly zero
- Linearity of variable : Independant variables are linearly dependant on the target variable


### Normality of Residuals

```
In [35]: y_pred = model.predict(X_test)
residuals = (y_test - y_pred)
sns.histplot(residuals, color='r')
plt.show()
```




From the histogram we can see that there is a negative skew in the distribution of residuals but it is close to a normal distribution

```
In [36]: stats.probplot(residuals, plot=plt)
plt.show()
```



### Test for Homoscedasticity

```
In [37]: plt.scatter(y_pred, residuals)  
plt.show()
```



There is no homoscedasticity present in the data, because the plot is not creating a cone type shape

### Insights -

- Multicollinearity is present in the data.
- The variables CGPA, GRE Score and TOEFL Score have a strong relationship with the target variable(Chance of Admit). These variables are also highly correlated to themselves.
- Mean of residuals is close to 0.
- Independent variables are linearly correlated with dependent variables.

### Recommendations -

- CGPA ,GRE Score, Toefl Score are important in making the prediction for Chance of Admit.
- All the exam scores are highly correlated. So it is recommende to add more independent features for better prediction.
- Independent variables like internships, projects completed can be added to the dataset to improve the prediction of the model.